

Faster than C?

Parsing binary data in JavaScript

@felixge



transloadit.com

Faster than C?

Sorry about the
“title bait”

High
performance
JavaScript

JavaScript vs C

Good vs Evil

Good Parts vs Evil

Bad Parts vs Evil


early 2010

No MySQL module
for node.js

early 2010

All we had was
NoSQL Libraries

early 2010



Yuichiro MASUI

masuidrive



JavaScript ★ 107 🔗 10

Pure JS / No C/C++


Before Buffers
became usable

The Parser was using
JavaScript Strings

Node.js Trivia

**“Buffers” used to be
called “Blobs”**

For 3 min and 15 sec



RIP Blobs

Sun Dec 13 08:39:20 2009

-

Sun Dec 13 08:42:45 2009

Anyway

mysql can be done
without libmysql



Felix Geisendörfer

felixge



node-mysql

A pure node.js JavaScript Client implementing the MySQL protocol.


Last updated 4 minutes ago

JavaScript

★ 1,364

🔗 161

No good deed goes
unpunished



Sir Isaac Newton

Third Law of Motion

“When a first body exerts a force F_1 on a second body, the second body simultaneously exerts a force $F_2 = -F_1$ on the first body. This means that F_1 and F_2 are equal in magnitude and opposite in direction.”

Third Law of Github

“When a first person pushes a library L1 into a remote repository, a second person simultaneously starts working on a second library L2 which will be equally awesome, but in a different way.”

<3 Github!



Oleg Efimov

Sannis



node-mysql-libmysqlclient

Asynchronous MySQL binding for Node.js


Last updated 25 days ago

JavaScript

★ 165

🔗 28





Of course.

libmysql = C

my library = JavaScript

C > JS, right?

But V8!

And Crankshaft!!

Node.js !!!1!

Was I
living a lie?


Kind of

V8 / Node

= Tools

Performance
is not a tool

Performance is
hard work &
data analysis



Third Law of Github



Brian White

mscdex



node-mariasql

A node.js binding to MariaDB's non-blocking (MySQL-compatible) client library


Last updated 4 days ago

C++

★ 23

🔗 0






Time to give up?



NEVER!

New Parser



Third law of Github?

Endgame

**Last bottleneck:
Creating JS Objects**



Also: MySQL Server saturated

Anyway

How to write fast JS

Does not work

Profiling

- Good for spotting small functions with stupid algorithms performing many iterations
- Bad for complex functions with many primitive operations

Taking performance advice from strangers

- Good for ideas & inspiration
- But useless when applied cargo-cult style

Does Work

BDD

~~Behavior Driven
Development~~

Benchmark Driven Development

Benchmark Driven Development

- Similar to test driven development
- Use it when performance is an explicit design goal
- Benchmark first > benchmark after !

```
1 function benchmark() {  
2     // intentionally empty  
3 }
```

```
1 while (true) {  
2     var start = Date.now();  
3     benchmark();  
4     var duration = Date.now() - start;  
5     console.log(duration);  
6 }
```

Benchmark Driven Development

- Next step: Implement a tiny part of your function
- Example: Parse headers of MySQL packets
- Look at impact, tweak code, repeat

Example Results

- try...catch is ok
- big switch statement = bad
- function calls = very cheap
- buffering is ok

Favorite

```
1 function parseRow(columns, parser) {  
2     var row = {};  
3     for (var i = 0; i < columns.length; i++) {  
4         row[columns[i].name] = parser.readColumnValue();  
5     }  
6     return row;  
7 }
```

Make it faster!

```
1 var code = 'return {\n';
2
3 columns.forEach(function(column) {
4     code += '"' + column.name + '":' + 'parser.readColumnValue(),\n';
5 });
6
7 code += '};\n';
8
9 var parseRow = new Function('columns', 'parser', code);
```

->

```
1 function parseRow(columns, parser) {  
2     return {  
3         id      : parser.readColumnValue(),  
4         title   : parser.readColumnValue(),  
5         body    : parser.readColumnValue(),  
6         created : parser.readColumnValue(),  
7         updated : parser.readColumnValue(),  
8     };  
9 }
```

eval = awesome


Data analysis


- Produce data points as tab separated values
- Add as many VM/OS metrics as you can get to every line
- Do not mix data and analysis !!

Recommended Tools

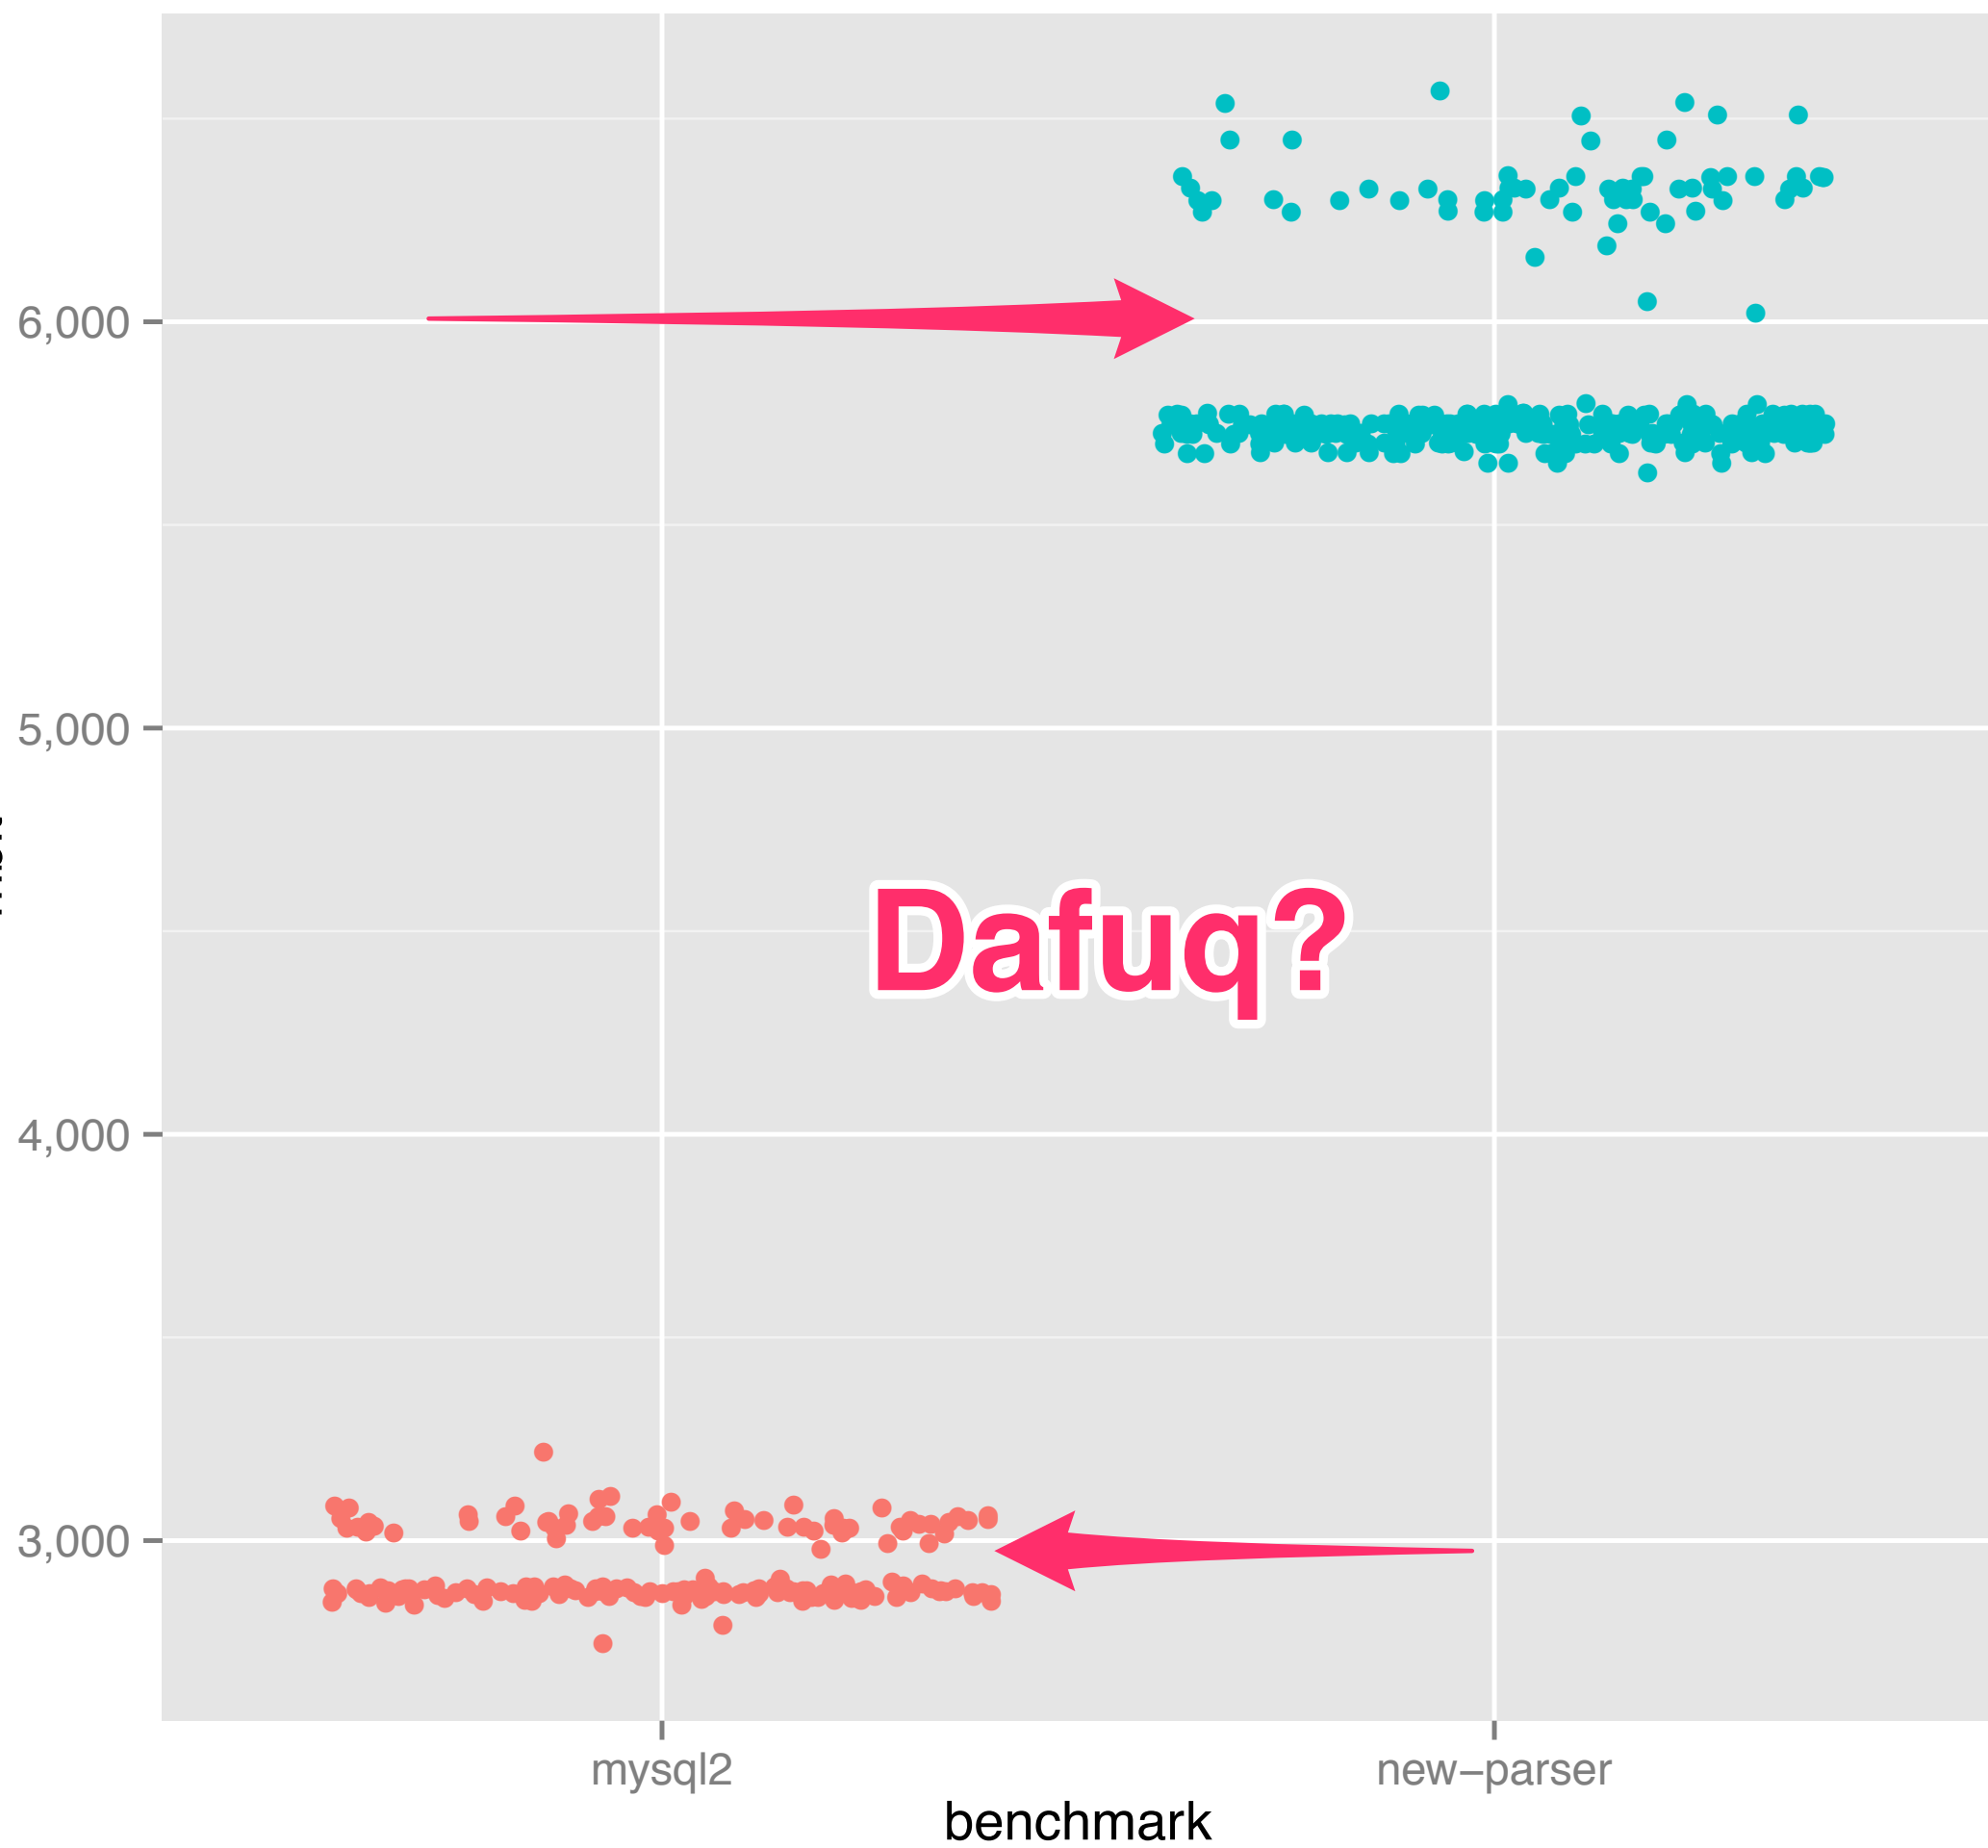
- `node benchmark.js | tee results.tsv`
- R Programming language (with ggplot2) !
- Makefiles, Image Magick, Skitch

Why?







mbit




benchmark

- mysql2
- new-parser







Collect data ->

Analyze it ->

Find problems ->

Tweak the code ->

Repeat

Thank you

Thank you !

github.com/felixge/faster-than-c

All benchmarks, results and analysis scripts